

Importing the Dependencies

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

Data Collection and Processing

```
#loading the csv data to a Pandas Dataframe
heart_data=pd.read_csv('/content/heart_disease_dataset.csv')
```

```
#print first five rows of the dataset
heart_data.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	sl
0	63	1	3	145	233	1	0	150	0	2.3	
1	37	1	2	130	250	0	1	187	0	3.5	
2	41	0	1	130	204	0	0	172	0	1.4	
3	56	1	1	120	236	0	1	178	0	0.8	
4	57	0	0	120	354	0	1	163	1	0.6	

Saved successfully!

×

```
#print last 5 rows of dataset
heart_data.tail()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	sl
298	57	0	0	140	241	0	1	123	1	0.2	
299	45	1	3	110	264	0	1	132	0	1.2	
300	68	1	0	144	193	1	1	141	0	3.4	
301	57	1	0	130	131	0	1	115	1	1.2	
302	57	0	1	130	236	0	0	174	0	0.0	

```
#number of rows and columns in the Dataset
heart_data.shape
```

(303, 14)

```
#getting some info about the data
heart_data.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302

```
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps     303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
#checking for missing values
heart_data.isnull().sum()
```

```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
thal     0
target   0
dtype: int64
```

Saved successfully!

```
#statistical measures about the data
heart_data.describe()
```

	age	sex	cp	trestbps	chol	fb
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.14851
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.35619
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.00000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.00000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.00000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.00000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.00000

```
#checking the distribution of Target Variable
```

```
heart_data['target'].value_counts()
```

```
1    165
0    138
Name: target, dtype: int64
```

1-->Defective Heart

0-->Healthy Heart

Splitting the Target and Features column

```
X=heart_data.drop(columns='target',axis=1)
```

```
Y=heart_data['target']
```

```
print(X)
```

	age	sex	cp	trestbps	chol	...	exang	oldpeak	slope	ca	thal
0	63	1	3	145	233	...	0	2.3	0	0	1
1	37	1	2	130	250	...	0	3.5	0	0	2
2	41	0	1	130	204	...	0	1.4	2	0	2
3	56	1	1	120	236	...	0	0.8	2	0	2
4	57	0	0	120	354	...	1	0.6	2	0	2
..
298	57	0	0	140	241	...	1	0.2	1	0	3
299	45	1	3	110	264	...	0	1.2	1	0	3
300	68	1	0	144	193	...	0	3.4	1	2	3
301	57	1	0	130	131	...	1	1.2	1	1	3
					236	...	0	0.0	1	1	2

Saved successfully!



[303 rows x 12 columns]

```
print(Y)
```

```
0    1
1    1
2    1
3    1
4    1
..
298  0
299  0
300  0
301  0
302  0
Name: target, Length: 303, dtype: int64
```

Splitting the Data into Training data and Testing data

```
X_train, X_test,Y_train, Y_test=train_test_split(X,Y, test_size=0.2, stratify=Y,random_state=2)
```

```
print(X_train.shape,X_test.shape)
```

(242, 13) (61, 13)

MODEL TRAINING

LOGISTIC REGRESSION

```
model=LogisticRegression()
```

```
#training the LogisticRegression model with Training Data  
model.fit(X_train, Y_train)
```

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)  
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
                    intercept_scaling=1, l1_ratio=None, max_iter=100,  
                    multi_class='auto', n_jobs=None, penalty='l2',  
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,  
                    warm_start=False)
```

MODEL EVALUATION

Saved successfully!



```
#accuracy on training data  
X_train_prediction=model.predict(X_train)  
training_data_accuracy=accuracy_score(X_train_prediction,Y_train)
```

```
print('Accuracy on Training data:',training_data_accuracy)
```

Accuracy on Training data: 0.8512396694214877

```
#accuracy on test data  
X_test_prediction=model.predict(X_test)  
test_data_accuracy=accuracy_score(X_test_prediction,Y_test)
```

```
print('Accuracy on Test data:',test_data_accuracy)
```

Accuracy on Test data: 0.819672131147541

BUILDING A PREDICTIVE SYSTEM

```
input_data = (66,0,3,150,226,0,1,114,0,2.6,0,0,2)
```

```
# change the input data to a numpy array
input_data_as_numpy_array= np.asarray(input_data)

# reshape the numpy array as we are predicting for only on instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)

prediction = model.predict(input_data_reshaped)
print(prediction)

if (prediction[0]== 0):
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')
```

```
[1]
The Person has Heart Disease
```

Saved successfully!



✓ 0s completed at 2:37 AM

